

Chapter 10: Azure AI Architecture

Learning Objectives

By the end of this chapter, you will:

- Understand the core Azure services used to build enterprise AI solutions.
- Learn common Azure AI reference architectures and deployment patterns.
- Develop the ability to discuss Azure AI architectures confidently with engineering teams and executives.

Beginner Explanation

Azure AI Architecture refers to the collection of Microsoft Azure services used to build, deploy, and operate AI applications.

Think of Azure as a city, where each service has a specific responsibility:

Azure Service	Analogy
Azure OpenAI	Brain
Azure AI Search	Library
Azure Functions	Workers
Azure Blob Storage	Warehouse
Azure Monitor	Security Cameras
Azure Key Vault	Safe
Azure Front Door	Main Entrance

Together, these services form an enterprise-grade AI platform.

Typical AI User Journey

1. User submits a request.
2. Application receives the request.
3. Azure OpenAI processes it.
4. Azure AI Search retrieves enterprise data.
5. Azure Functions execute workflows.
6. Response is returned to the user.

Most enterprise AI applications follow this pattern.

Intermediate Explanation

Enterprise Azure AI architectures generally follow one of three patterns.

Pattern 1: Simple LLM

```text id="7cth0w" User → Application → Azure OpenAI → Response

Best for:

- Content generation
- Summarization
- Translation

---

#### ### Pattern 2: RAG Architecture

```text id="50ivmr"

User → Application

↓

Azure AI Search

↓

Azure OpenAI

↓

Response

Best for:

- Knowledge assistants
- Policy bots
- Internal search solutions

Pattern 3: AI Agent Architecture

```text id="b64vhy" User ↓ AI Agent ↓ Enterprise Tools ↓ Azure OpenAI ↓ Actions

Best for:

- Autonomous workflows
- Incident management
- Program management assistants

---

#### ### Azure AI Design Principles

| Principle         | Description                  |
|-------------------|------------------------------|
| Scalability       | Handle growing demand        |
| Reliability       | Maintain availability        |
| Security          | Protect enterprise data      |
| Observability     | Monitor system health        |
| Cost Optimization | Manage cloud spend           |
| Compliance        | Meet regulatory requirements |

---

## ## Why This Matters for an AI TPM

| Responsibility         | Impact                                                               |
|------------------------|----------------------------------------------------------------------|
| Delivery               | Ensures AI solutions are designed and deployed successfully.         |
| Risk                   | Identifies architectural risks before production deployment.         |
| Stakeholder Management | Communicates architecture decisions to technical and business teams. |
| Budget                 | Evaluates cost implications of Azure services and scaling decisions. |

---

## ## Enterprise Example

### ### Business Problem

A multinational telecommunications company wanted to deploy an AI-powered support assistant for more than 50,000 employees across 20 countries.

#### Requirements:

- Multi-region deployment
- Enterprise security
- Less than 3-second response time
- Integration with Microsoft Teams

### ### Solution

Implement a scalable Azure AI reference architecture.

### ### Architecture

- Azure OpenAI
- Azure AI Search
- Azure Functions
- Azure Front Door
- Azure API Management
- Azure Monitor
- Azure Key Vault

- Microsoft Teams

### ### Outcome

| Metric                   | Result      |
|--------------------------|-------------|
| Availability             | 99.96%      |
| Employee Adoption        | 88%         |
| Response Time            | 2.1 Seconds |
| Support Ticket Reduction | 57%         |

### ### Lessons Learned

- Azure Front Door improves global performance.
- AI Search significantly improves answer quality.
- Multi-region deployments increase resilience.
- Governance should be established before scaling.

---

## ## Azure Implementation

### ### Services Used

| Service          | Purpose                           |
|------------------|-----------------------------------|
| Azure OpenAI     | LLM inference                     |
| Azure AI Search  | Enterprise retrieval              |
| Azure Functions  | Workflow execution                |
| Azure AI Foundry | Prompt management and evaluations |

### ### Supporting Services

| Service              | Purpose                |
|----------------------|------------------------|
| Azure Front Door     | Global routing         |
| API Management       | API gateway            |
| Azure Monitor        | Observability          |
| Key Vault            | Secrets management     |
| Blob Storage         | Document storage       |
| Application Insights | Performance monitoring |

---

### ### Implementation Workflow

#### #### Step 1: Define Architecture Requirements

Document:

- Users

- Traffic estimates
- Security needs
- Compliance requirements
- Performance targets

---

#### #### Step 2: Select Reference Architecture

Options:

- Simple LLM
- RAG
- Multi-Agent
- Hybrid

---

#### #### Step 3: Configure Security

Implement:

- Private Endpoints
- RBAC
- Key Vault
- Network Security Groups

---

#### #### Step 4: Deploy Infrastructure

Provision:

- Azure OpenAI
- AI Search
- Functions
- Monitoring services

---

#### #### Step 5: Enable Observability

Monitor:

- Availability
- Latency
- Token usage
- Cost
- Errors

---

```

Mermaid Diagram

```mermaid id="6q3jpv"
graph TD
  A[Users]
  A --> B[Azure Front Door]
  B --> C[API Management]
  C --> D[Azure OpenAI]
  D --> E[Azure AI Search]
  E --> F[Azure Functions]
  F --> G[Blob Storage]
  G --> H[Azure Monitor]
  H --> I[Operations Team]

```

Enterprise RAG Architecture

```

```mermaid id="u2a8mt" graph TD
 A[Enterprise Documents]
 A --> B[Blob Storage]
 B --> C[Embeddings]
 C --> D[Azure AI Search]
 D --> E[Azure OpenAI]
 E --> F[Microsoft Teams]
 F --> G[Employees]

```

---

```

Multi-Agent Architecture

```mermaid id="s8v8qa"

```

```

graph LR
  A[User]
  A --> B[Coordinator Agent]
  B --> C[HR Agent]
  B --> D[IT Agent]
  B --> E[TPM Agent]
  B --> F[Finance Agent]
  C --> G[Enterprise Systems]
  D --> G
  E --> G
  F --> G

```

TPM Checklist

- ☐ Define architecture requirements.
- ☐ Identify stakeholders.
- ☐ Estimate monthly costs.
- ☐ Conduct architecture reviews.
- ☐ Implement security controls.
- ☐ Configure monitoring.
- ☐ Conduct performance testing.
- ☐ Review disaster recovery plans.
- ☐ Validate compliance requirements.
- ☐ Obtain architecture approvals.

Templates

Azure Architecture Review Template

Item	Status
Security Approved	Yes/No
Monitoring Configured	Yes/No
DR Plan Completed	Yes/No
Costs Estimated	Yes/No

Item	Status
Compliance Approved	Yes/No

Monthly Cost Tracker

Service	Estimated Cost
Azure OpenAI	
Azure AI Search	
Azure Functions	
Monitoring	
Total	

Risk Register

Risk	Impact	Mitigation
Regional Outage	High	Multi-region deployment
Cost Overruns	Medium	Budget alerts
Data Leakage	High	RBAC and encryption
High Latency	Medium	Caching and optimization

Common Mistakes

1. Selecting the wrong architecture pattern.
2. Ignoring security requirements.
3. Underestimating operational costs.
4. Skipping monitoring implementation.
5. Neglecting disaster recovery planning.
6. Deploying without performance testing.
7. Failing to involve stakeholders early.

Key Takeaways

- Azure provides a complete ecosystem for building enterprise AI applications.
- RAG and AI Agent architectures are becoming the standard for enterprise deployments.
- Security, scalability, and observability should be considered from the beginning.
- Azure Front Door, API Management, and Monitoring services are critical components of production systems.

- AI TPMs must understand Azure reference architectures to effectively lead AI programs.
- Architecture decisions have a direct impact on cost, reliability, and business outcomes.

Footer

charan.ai | Page 49